

The Distributed Intelligence of Life

*How AI Must Learn from Biology Without Destroying What
Makes Biology Intelligent*

By Sage Clokey

A technical companion to "The Living Republic of Conscience"

I. The Central Problem

Modern artificial intelligence, when applied to biological data, commits the same error that central planning commits when applied to economies. It aggregates.

It computes mean expression levels across tissues. It builds consensus sequences across populations. It reports average fitness across environments. It measures error rates across experiments and calls the residual noise. It reduces millions of individual data points — each one the record of a specific living agent acting in a specific context — into summary statistics that no individual agent would recognize as its own.

This is not a minor methodological preference. It is a fundamental misunderstanding of where biological information lives.

The information does not live in the average. It lives in the difference.

The single nucleotide polymorphism that distinguishes a functional protein from a misfolded one. The one cell in a tissue of billions that escapes immune surveillance and becomes cancerous. The specific transposon insertion that rewires a regulatory network and enables a novel phenotype. The particular epigenetic state that causes two genetically identical cells to express entirely different gene programs. The precise context — genomic, cellular, environmental — that determines whether a mutation is beneficial, neutral, or lethal.

These are not deviations from the signal. They are the signal. When we average across them, we do not clarify the data. We destroy it.

II. The Biological Knowledge Problem

Friedrich Hayek demonstrated in 1945 that the knowledge required for rational economic coordination is dispersed among millions of individuals, exists in forms that cannot be articulated or transmitted to a central authority, and changes constantly. No central planner can possess this knowledge. Therefore, no central plan can substitute for the distributed process — the price system — that uses it.

James Shapiro demonstrated that the same structure governs biological systems. Every cell possesses local knowledge — its

position in the tissue, its metabolic state, its signaling environment, its epigenetic configuration, its history of encounters with stress and pathogens — that exists in no other cell. This knowledge is not merely difficult to centralize. It is impossible to centralize, because it is constituted by the cell's particular relationship to its particular context. Remove the context and the knowledge ceases to exist.

When a computational pipeline takes single-cell RNA sequencing data — data that captures the expression state of individual cells — and immediately clusters those cells into types, computes mean expression profiles for each cluster, and discards the individual measurements, it is performing the computational equivalent of Soviet central planning. It is replacing distributed knowledge with aggregate statistics and calling the result understanding.

The aggregate tells you that cells of type X express gene Y at level Z on average. It does not tell you that cell number 4,713 in cluster X has an expression profile that deviates from the mean in a pattern that, if you had preserved it, would have revealed the early signature of a state transition — a cell beginning to differentiate, or beginning to lose contact inhibition, or beginning to respond to a signal that its neighbors have not yet received.

That cell is the entrepreneur in Hayek's economy. It possesses knowledge that the aggregate cannot contain. It is acting on information that the summary statistic has destroyed. And the system's future behavior — whether the tissue remains healthy or develops a tumor, whether the organism adapts to a new environment or fails — depends on what that individual cell does with its individual knowledge.

Averaging is not analysis. It is the destruction of analysis.

III. Where the Data Lives

In economics, the data lives in prices — specific prices for specific goods in specific markets at specific times. The price of wheat in Chicago on a Tuesday morning is not the same datum as the average price of grain globally. The specific price carries information about local supply, local demand, local weather, local transportation costs, local expectations about the future. The average carries almost none of this.

In biology, the data lives in differences — specific differences between specific sequences, specific cells, specific organisms, specific contexts.

Between sequences. Two proteins may share 98% sequence identity. The 2% that differs determines whether one binds a receptor and the other does not, whether one folds in milliseconds and the other aggregates, whether one catalyzes a reaction at a rate sufficient for life and the other does not. The functional information is not in the 98% that is shared. It is in the 2% that diverges. A model trained to predict the consensus learns the background. A model trained to identify and interpret the differences learns the biology.

Between cells. In a tissue of ten thousand cells, nine thousand nine hundred and ninety-nine may be in a stable, differentiated state. One may be transitioning. That one cell — its specific expression profile, its specific chromatin state, its specific position in the tissue — carries more information about the system's dynamics than all the others combined. It is the leading edge of a process that the bulk measurement cannot see. A pipeline that averages across all

ten thousand cells does not merely fail to find this cell. It actively conceals it.

Between conditions. The same gene knocked out in two different genetic backgrounds may produce opposite phenotypes. The same mutation in two different environments may be beneficial in one and lethal in the other. The same protein expressed at two different levels may activate a pathway at low concentration and inhibit it at high concentration. The information is not in the gene, the mutation, or the protein considered in isolation. It is in the specific interaction between the element and its context. A model that learns "gene X does Y" without encoding the context has learned a falsehood that happens to be true on average.

Between timepoints. A cell's trajectory through state space — the specific sequence of expression changes it undergoes as it differentiates, responds to stress, or transitions to disease — carries information that no single snapshot can contain. Two cells may occupy the same state at a given moment but be moving in opposite directions. One is differentiating normally. The other is dedifferentiating toward malignancy. The static measurement cannot distinguish them. Only the trajectory — the difference between successive states — reveals what is happening.

The data is always in the difference. The difference is always specific. The specific is always contextual. And context is always local.

IV. How Current AI Destroys Biological Information

The standard machine learning pipeline for biological data follows a predictable sequence of information destruction:

Step 1: Aggregation. Raw measurements from individual cells, individual organisms, or individual experiments are pooled into summary statistics — means, medians, variance estimates, principal components. The justification is noise reduction. The effect is signal destruction.

Step 2: Dimensionality reduction. High-dimensional data — thousands of genes measured per cell, millions of variants per genome — is projected into low-dimensional representations. Principal component analysis, t-SNE, UMAP. The justification is visualization and tractability. The effect is the loss of precisely those dimensions that encode rare, context-dependent, or combinatorial signals.

Step 3: Classification. Continuous, context-dependent biological states are forced into discrete categories — cell type A, cell type B; healthy, diseased; functional, nonfunctional. The justification is interpretability. The effect is the erasure of boundary states, transitional states, and mixed states — exactly the states that carry the most dynamic information.

Step 4: Error minimization. The model is trained to minimize average error across the dataset. The loss function treats every prediction error as equally unimportant. A model that perfectly predicts the behavior of the 99% majority and completely fails on the 1% minority receives a score of 99% — and misses the only data that matters. The rare variant, the outlier cell, the context-

dependent exception — these are treated as noise to be minimized rather than signal to be understood.

Step 5: Consensus output. The trained model produces predictions that represent the central tendency of its training data. Asked to generate a protein sequence for a given function, it generates the average of all sequences that perform that function — a sequence that may not fold, may not function, and certainly does not represent any specific, tested, viable solution. It has learned the blur, not the grammar.

At every step, the pipeline moves further from the distributed, contextual, difference-rich reality of biology and closer to the aggregated, decontextualized, averaged fiction of central planning.

V. The Architecture of Distributed Intelligence

If we take seriously the claim that biological information is distributed, contextual, and encoded in differences, then the AI systems we build to learn from biology must be designed to preserve these properties rather than destroy them. This is not a matter of better algorithms applied to the same framework. It is a different framework entirely.

Principle 1: Learn from Differences, Not Summaries

The fundamental unit of biological learning should be the comparison — not the data point in isolation.

Instead of training a model on individual protein sequences labeled "functional" or "nonfunctional," train it on pairs: this sequence

functions, that sequence does not, and here is what differs between them. The model learns not "what a functional protein looks like on average" but "what specific changes convert a functional protein into a nonfunctional one, and vice versa."

This is contrastive learning applied with biological intention. The contrastive loss function does not ask "how close is this prediction to the population mean?" It asks "can you identify which specific differences between these two inputs account for their different outcomes?" The model is rewarded for precision about particulars, not accuracy about averages.

Protein language models like ESM-3 already move in this direction by learning from individual sequences positioned within an evolutionary context — each sequence is understood in relation to the others in its family. But the next generation of models should make the comparison explicit: learn the grammar of mutation, not just the grammar of sequence. Learn what changes do, not just what sequences are.

Principle 2: Preserve Individual Identity

Every cell is an agent. Every sequence is a sentence. Every organism is a citizen of its ecosystem. The AI system must preserve the identity of each individual datum as long as computationally possible, resisting the urge to aggregate.

In practice, this means:

Single-cell data should remain single-cell. Graph neural networks can model cell-cell interaction networks where each node retains its individual expression profile, its spatial position, and its signaling relationships with specific neighbors. The model learns

from the network structure without collapsing individual cells into cluster averages. Each cell remains a citizen of the tissue republic, with its own law and its own local knowledge.

Sequence data should remain sequence-level. Instead of building position-weight matrices that average across an alignment, use attention mechanisms that can query individual positions in individual sequences. The transformer architecture is naturally suited to this — its self-attention mechanism allows each position to attend to every other position, preserving the full combinatorial structure of the sequence rather than compressing it into a consensus.

Temporal data should remain trajectory-level. Instead of comparing snapshots at time A and time B, model the continuous trajectory of each individual cell or organism through state space. Recurrent architectures and neural ordinary differential equations can learn dynamics from individual trajectories without requiring that all individuals follow the same path.

The computational cost of preserving individual identity is real. It is also the price of not destroying the data. Central planning is computationally cheaper than a free market — you only need to optimize one objective function instead of coordinating millions of agents. But the efficiency is illusory, because the central plan optimizes the wrong function. The same is true in computational biology: the aggregated pipeline is faster, but it optimizes for accuracy on the average case at the cost of blindness to every case that matters.

Principle 3: Encode Context as a First-Class Input

The same mutation in different contexts produces different outcomes. Therefore, context is not a confounding variable to be controlled for. It is a primary variable to be learned from.

This means that every biological prediction should be conditioned on context:

- **Genomic context.** What is the surrounding sequence? What regulatory elements are nearby? What is the chromatin state? What other variants are present in the same genome?
- **Cellular context.** What cell type? What tissue? What developmental stage? What metabolic state? What signaling environment?
- **Environmental context.** What temperature? What nutrient availability? What stressors are present? What other organisms are in the community?
- **Evolutionary context.** What lineage does this sequence belong to? What selection pressures has it experienced? What other solutions to the same problem exist in related lineages?

A model that predicts "this mutation is deleterious" without specifying the context has made a statement that is either trivially true (on average, most mutations are deleterious) or dangerously false (this specific mutation, in this specific context, may be the adaptive innovation that saves the lineage). Context-free prediction is the biological equivalent of price-free allocation — technically possible, practically useless, and systematically misleading.

Architecturally, this means building models that take context as an explicit input — not as a batch label to be controlled for in post-

processing, but as a structured representation that the model learns to integrate with the primary data. Conditional generation, context-aware attention, and multi-modal architectures that jointly model sequence, structure, expression, and environment are steps in this direction.

Principle 4: Reward Discrimination, Not Consensus

The loss function defines what the model learns. If the loss function rewards closeness to the mean, the model learns the mean. If it rewards the ability to discriminate between specific cases, the model learns what makes specific cases different.

For biological AI, the loss function should:

- **Penalize false negatives on rare events more than false positives on common events.** The rare cell, the rare variant, the rare phenotype — these carry disproportionate information. A model that detects 100% of common cell types and 0% of rare transitional states has learned nothing useful about the system's dynamics.
- **Reward correct identification of causal differences.** Given two sequences with different functions, can the model identify which residue changes are responsible? Given two cells with different fates, can the model identify which expression differences are predictive? The model should be evaluated on its ability to localize the specific differences that matter, not on its average accuracy across all positions.
- **Penalize overconfident consensus predictions.** A model that outputs "this sequence will fold with 99% confidence" when trained on averaged data should be penalized for false

certainty. Biological systems are inherently variable, and the model's uncertainty should reflect the genuine heterogeneity of the underlying population — not collapse it into a single-point estimate.

Principle 5: Generate Specific Instances, Not Averages

When a generative model is asked to produce a protein sequence for a given function, it should produce a specific, viable sequence — not the average of all sequences that perform that function.

The average sequence is a statistical ghost. It exists in no organism. It has been tested by no environment. It has survived no selection. It is the biological equivalent of a centrally planned price — a number that satisfies a mathematical criterion but corresponds to no real transaction.

A generative model for biology should operate like natural genetic engineering itself: it should produce specific variants, situated in specific contexts, to be tested by specific environments. It should be capable of generating diverse solutions to the same problem — multiple different sequences that achieve the same function through different structural strategies — because this is what biology does. Evolution does not converge on one answer. It explores a space of viable answers, each adapted to its particular context.

This means training generative models not just on the distribution of functional sequences but on the diversity within that distribution. The model should learn that there are many ways to be a kinase, many ways to bind DNA, many ways to catalyze a reaction — and it should be able to sample from the full range of solutions rather than collapsing to the mode.

Principle 6: Distributed Verification

No prediction should be accepted on the authority of the model alone. Every output of a biological AI system must be verified through the same distributed process that governs biological creation: experimental testing in real environments with real organisms.

This is not merely a methodological precaution. It is an architectural principle. The AI system should be designed as one agent in a distributed network — a participant in the cycle of hypothesis, experiment, and revision, not an oracle that issues final pronouncements. Its predictions are proposals, not conclusions. Its generated sequences are candidates, not solutions. Reality is the only validator.

The temptation to treat AI as a central planner — to trust its predictions because they emerge from a model trained on vast data — is the same temptation that Hayek identified in economics: the fatal conceit that because we can observe the order, we can command it. We cannot. The order is too complex, too distributed, too alive. The best we can do is propose, test, revise, and propose again — the same distributed, directed trial and error that life has used for four billion years.

VI. The English-to-Protein Bridge

The Living Age project aims to build an English-to-Protein bridge: the capacity to describe a desired biological function in natural

language and translate that description into a viable protein sequence using AI tools like ESM-3 and AlphaFold.

This bridge, if built correctly, is an act of literacy — learning to speak the language that life already speaks. If built incorrectly, it is an act of central planning — imposing human designs on living systems without regard for the distributed knowledge that makes those systems work.

The principles above define what "built correctly" means:

The bridge must learn from differences. It must understand not just "what is a kinase" but "what makes this kinase different from that kinase, and why does that difference matter in this context." The training data must preserve the variation within functional families, not collapse it into consensus profiles.

The bridge must preserve context. A request for "a protein that binds target X" is incomplete without specifying: in what organism? In what cellular compartment? At what concentration? In the presence of what other proteins? Under what environmental conditions? The bridge must require context as input and condition its outputs on that context.

The bridge must generate specific, viable sequences. Not the average protein that binds target X, but a particular protein — with a particular structure, a particular binding mode, a particular set of properties — that can be synthesized, expressed, and tested in a real system.

The bridge must participate in distributed verification. Its outputs must be treated as hypotheses to be tested, not as solutions to be deployed. The cycle of prediction, synthesis, testing, and

revision is not a limitation of the system. It is the system. It is natural genetic engineering conducted with human tools.

The bridge must respect the grammar without claiming to command it. The deep rules of protein folding, gene regulation, and cellular cooperation were not designed by any human mind. They were discovered through four billion years of distributed experimentation by trillions of living agents. The AI model has learned to read some of this grammar. It has not learned to replace it. Humility before the living order is not a sentimental posture. It is a technical requirement. The system is more complex than any model we can build, and acting as though it were not is the surest path to failure.

VII. From Central Planning to Distributed Intelligence

The history of artificial intelligence in biology has largely been a history of central planning applied to distributed systems. We have built models that aggregate, classify, average, and predict — models that treat biological data the way a Soviet ministry treats economic data: as raw material to be processed into summary reports for decision-makers who sit above the system and issue commands downward.

This approach has produced real results. Bulk RNA-seq analysis has identified disease biomarkers. Genome-wide association studies have linked variants to phenotypes. AlphaFold has predicted protein structures with remarkable accuracy. These achievements are genuine and should not be dismissed.

But they are achievements of the kind that central planning can produce: accurate descriptions of the average case, useful for problems where the average case is what matters. They fail — systematically and predictably — for problems where the individual case is what matters. For understanding why this patient responds to a drug and that patient does not. For predicting how this specific mutation will behave in this specific genome. For designing a protein that functions not in the abstract but in a particular cellular environment.

These are the problems that matter most, and they are the problems that require a different approach — an approach that treats biological data the way a free market treats economic data: as the distributed knowledge of millions of agents, each possessing information that cannot be centralized, each acting in a context that cannot be averaged away.

The technical challenge is real. Preserving individual-level data is computationally expensive. Learning from differences rather than summaries requires more sophisticated architectures. Encoding context as a first-class input multiplies the dimensionality of the problem. Generating diverse, specific outputs rather than consensus predictions demands generative models of greater capacity and subtlety.

But these are engineering challenges, not fundamental obstacles. The fundamental obstacle is conceptual: the assumption, inherited from the statistical tradition, that the purpose of data analysis is to extract signal from noise by averaging across individuals. In biology, the individuals are the signal. The average is the noise.

The Living Age proposes a different relationship between artificial intelligence and living systems — one modeled not on central

planning but on distributed intelligence. AI as a participant in the republic of life, not a planner above it. AI as a reader of grammar, not a commander of syntax. AI as one agent among many in the distributed, directed trial and error that has always been the engine of biological creation.

The data is in the difference. The knowledge is in the individual. The intelligence is in the distribution. Build accordingly.

This essay is a technical companion to "The Living Republic of Conscience: Austrian Economics, Natural Genetic Engineering, and the Design Rules of Life." Together, they describe the philosophical foundation and the technical architecture of the Living Age — a vision of human-biological collaboration grounded in distributed intelligence, individual agency, and respect for the living order.