

CHAPTER 2

Bioethics and Moral Theories

As we have seen, the moral life is dynamic, complex, and inescapable. In it we wrestle with momentous questions of moral value and moral rightness. We assert, challenge, accept, and reject moral statements. We make moral judgments about the rightness of actions, the goodness of persons or their character, and the moral quality and worth of our lives. Through general moral norms or principles, we direct our actions and inform our choices. We formulate and critique moral arguments, thereby testing what we know or think we know about moral realities. We do all this and one thing more: We naturally and unavoidably venture into the realm of moral theory, trying to see the larger moral meaning behind particular situations and precepts. In this chapter, we explore this realm and try to discern how it fits into the moral life in general and into bioethics in particular.

THE NATURE OF MORAL THEORIES

In science, theories help us understand the empirical world by explaining the causes of events, why things are the way they are. The germ theory of disease explains how particular diseases arise and spread in a human population. The heliocentric (sun-centered) theory of planetary motion explains why the planets in our solar system behave the way they do. In ethics, moral theories have a similar explanatory role. A **moral theory** explains not why one event causes another but why an action is right or wrong or why a person or a person's character is good or bad. A moral theory tells us what it is about an action that *makes it right*, or what it is about a person that *makes him or her good*. The divine command

theory of morality, for example, says that right actions are those commanded or willed by God. Traditional utilitarianism says that right actions are those that produce the greatest happiness for all concerned. These and other moral theories are attempts to define rightness or goodness. In this way, they are both more general and more basic than moral principles or other general norms.

Moral theorizing—that is, making, using, or assessing moral theories or parts of theories—is normal and pervasive in the moral life, though it is often done without much recognition that theory is playing a part in the deliberations. Whenever we try to understand what a moral property such as rightness or goodness means, or justify a moral principle or other norm, or resolve a conflict between two credible principles, or explain why a particular action or practice is right or wrong, or evaluate the plausibility of specific moral intuitions or assumptions, we do moral theorizing. In fact, we *must* theorize if we are to make headway in such investigations. We must stand back from the situation at hand and try to grasp the larger pattern that only theory can reveal.

Moral theories that concentrate on right and wrong actions are known as theories of obligation (or duty) or simply as *theories of right action*. The divine command theory and utilitarianism are theories of right action. Philosophers often distinguish these from moral theories that focus on good and bad persons or character—so-called *virtue-based theories*. Virtue ethics (covered later in this chapter) is a prime example.

How do moral theories fit into our everyday moral reasoning? In answering that, let's focus on theories of right action, probably the most

influential type in bioethics. First, moral theories can figure directly in our moral arguments. As we saw earlier, moral arguments contain both moral and nonmoral premises. A moral premise can consist of a moral principle, a moral rule (a less general norm derived from or based on a principle), or a claim expressing a central tenet of a moral theory. Using such a tenet, someone might argue, for example, that stem-cell research should be fully funded rather than halted altogether because such a step would eventually lead to a greater benefit for more people, and right actions (according to utilitarianism) are those that result in the greatest overall benefit for the greatest number. Thus the fundamental moral standard of utilitarianism becomes a premise in an argument for a specific action in a particular case.

Second, theories can have an indirect impact on moral arguments because principles appealed to are often supported in turn by a moral theory. The principles can be either derived from or supported by the theory's account of right and wrong action. Consider the prohibition against murder, the basic precept that it is wrong to take the life of an innocent person. This principle can be drawn from theories built around the fundamental notion of respect for persons. As one such theory would have it, murder is wrong because it treats people not as persons with inherent worth but as mere things to be used or dispensed with as one wishes.

Some people are tempted to deduce from all this that moral theories are the dominant force in moral reasoning as well as in the moral life. This view would be an oversimplification. By design, moral theories are certainly more general in scope than moral principles, rules, or judgments. But from this fact it does not follow that theories alone are the ultimate authority in moral deliberations. For one thing, to be truly useful, moral theories must be filled out with details about how to apply them in real life and the kinds of cases to which they are relevant. For another, there is more to morality than what can be captured in the general norms of a theory.

There is also the testimony of the particular, the evidence of individual moral judgments.

Our moral deliberations, then, involve both the general and the particular. Suppose we embrace a moral theory that seems to offer us a plausible explanation of what makes an action right or wrong. When we must decide which action is morally right in a particular situation, we look to our theory for general guidance. From our theory we may glean a set of moral principles that seem to apply to the case at hand. If the principles lead us to conflicting choices, we look again to the theory for insight in resolving the conflict. But we also must take into account our considered judgments about the case. (We may also formulate considered judgments about the relevant principles or rules.) If our considered judgments and the deliverances of our theory are consistent with one another, we have additional assurance that our decision in the case is correct. If our judgments clash with our theory or principles, we must decide which to revise or discard—for critical reasoning demands that our beliefs be coherent, that they do not harbor contradictions. If we believe our judgments to be more credible than the implications of our theory (or principles), we may modify the theory accordingly (or, rarely, regard the theory as irreparable and give it up). But if the theory seems more credible in this case, we may conclude that our judgment is untrustworthy and set it aside.

So a moral theory can show us what is important and reasonable in morality, guiding our judgments through overarching insights that may help us with specific cases and issues, sometimes correcting erring judgments along the way. Our considered judgments are fallible indicators of moral common sense and are checks against wayward theory or flawed principle. In bioethics, both of these moral resources are highly respected and widely used.

INFLUENTIAL MORAL THEORIES

Several moral theories have played major roles in bioethics, and they continue to influence how people think about bioethical issues. Theories of

right action (in contrast to virtue-based theories) have dominated the field, each usually based on one of two broad views about the essential character of right actions. **Consequentialist** moral theories insist that the rightness of actions depends solely on their consequences or results. The key question is what or how much good the actions produce, however *good* is defined. **Deontological** (or **nonconsequentialist**) theories say that the rightness of actions is determined not solely by their consequences but partly or entirely by their intrinsic nature. For some or all actions, rightness depends on the kind of actions they are, not on how much good they produce. A consequentialist theory, then, may say that stealing is wrong because it causes more harm than good. But a deontological theory may contend that stealing is inherently wrong regardless of its consequences, good or bad.

Utilitarianism

The leading consequentialist theory is **utilitarianism**, the view that right actions are those that result in the most beneficial balance of good over bad consequences for everyone involved. It says we should maximize the nonmoral good (the *utility*) of everyone affected, regardless of the contrary urgings of moral rules or unbending moral principles. Various forms of utilitarianism differ in how they define utility, with some equating it with happiness or pleasure (the hedonistic view), others with satisfaction of preferences or desires or some other intrinsically valuable things or states such as knowledge or perfection.

In applying the utilitarian moral standard (the greatest good, everyone considered), some moral philosophers concentrate on specific acts and some on rules covering kinds of acts. The former approach is called **act-utilitarianism**, the idea that the rightness of actions depends solely on the relative good produced by *individual actions*. An act is right if in a particular situation it produces a greater balance of good over bad than any alternative acts; determining rightness is a matter of weighing the effects of each possible act. The latter approach, known as **rule-utilitarianism**,

avoids judging rightness by specific acts and focuses instead on *rules governing categories of acts*. It says a right action is one that conforms to a rule that, if followed consistently, would create for everyone involved the most beneficial balance of good over bad. We are to adhere to the rules because they maximize the good for everyone considered—even though a given act may produce bad effects in a particular situation.

The classic version of utilitarianism was devised by English philosopher Jeremy Bentham (1748–1832) and given more detail and plausibility by another English philosopher, John Stuart Mill (1806–1873). Classic utilitarianism is hedonistic in that the utility to be maximized is pleasure, broadly termed happiness, the only intrinsic good. A right action produces more net happiness (amounts of happiness minus unhappiness) than any alternative action, everyone considered. As Mill put it,

[Actions] are right in proportion as they tend to promote happiness, wrong as they tend to produce the reverse of happiness. By “happiness” is intended pleasure, and the absence of pain; by “unhappiness,” pain and the privation of pleasure.¹

Bentham and Mill, however, had different ideas about what happiness entailed, as do many philosophers today. Bentham thinks that happiness is one-dimensional: It is pleasure, pure and simple, something that varies only in the amount that an agent can experience. On this scheme, it seems that the moral ideal would be to experience maximum amounts of pleasure, as does the glutton or the debauchee. But Mill thinks that pleasures can vary in quality as well as quantity. For him, there are lower and higher pleasures—the lower and inferior ones indulged in by the glutton and his ilk and the higher and more satisfying ones found in such experiences as the search for knowledge and the appreciation of art and music. Mill famously sums up this contrast by saying, “It is better to be a human being dissatisfied than a pig satisfied; better to be Socrates dissatisfied than a fool satisfied.”²

Like all forms of utilitarianism, the classic formulation demands a strong sense of impartiality. When promoting happiness, we must not only take into account the happiness of everyone affected but also give everyone's needs or interests equal weight. Mill explains:

[The] happiness which forms the utilitarian standard of what is right conduct, is not the agent's own happiness, but that of all concerned. As between his own happiness and that of others, utilitarianism requires him to be as strictly impartial as a disinterested and benevolent spectator.³

In classic utilitarianism, the emphasis is on maximizing the total quantity of net happiness, not ensuring that it is rationed in any particular amounts among the people involved. This means that an action resulting in 1,000 units of happiness for 10 people is better than an action yielding only 900 units of happiness for those same 10 people—regardless of how the units of happiness are distributed among them. Classic utilitarians do want to allocate the total amount of happiness among as many people as possible (thus their motto, "the greatest happiness for the greatest number"). But maximizing total happiness is the fundamental concern whether everyone gets an equal portion or one person gets the lion's share.

How might utilitarianism apply to a bioethical issue? Consider this scenario: Johnny is a 10-year-old boy with cerebral palsy, emaciated and bed-ridden, hooked to feeding tubes and monitors, his body twisted in pain that is almost impossible to control, his days measured out by one agonizing surgical operation after another, locked in the mental life of an infant and acknowledged by all the experts to be without hope. His anguished parents, wanting desperately to end his suffering, beg the physician to give Johnny a lethal injection. What should the physician do?

Suppose in this case there are only two options: indefinitely maintaining Johnny in his present condition or carrying out the parents' wishes. An act-utilitarian might reason like this. Allowing the current situation to continue would

cause enormous unhappiness—Johnny's own physical agony, the unimaginable misery of the distraught parents, the anxiety of other family members and friends, and the distress and frustration of the physician and nurses who can do little more than stand by as Johnny withers away. On the other hand, administering the lethal injection would immediately end Johnny's pain and prevent future suffering. The parents would grieve for Johnny but would at least find some relief—and perhaps peace—in knowing that his torture was over. The medical staff would probably also be relieved for the same reason. There would, of course, also be possible negative consequences to take into account. In administering the lethal injection, the physician would be risking both professional censure and criminal prosecution. If her actions were to become public, people might begin to mistrust physicians who treat severely impaired children, undermining the whole medical profession. Perhaps the physician's action would lead to a general devaluing of the lives of disabled or elderly people everywhere. These dire consequences, however, would probably not be very likely if the physician acted discreetly. On balance, the act-utilitarian might say, greater net happiness (the least unhappiness) would result from the mercy killing, which would therefore be the morally permissible course.

A rule-utilitarian might judge the situation differently. The key question would be which rule if consistently followed would produce the greatest net happiness. Let us say that there are only two rules to consider. One says "Do not kill seriously impaired children, regardless of their suffering or the wishes of their parents." The other one is "Killing seriously impaired children is permissible if they are suffering severely and improvement is hopeless." The rule-utilitarian might reason that consistently following the second rule would have terrible consequences. It would cause widespread suspicion about the actions and motives of physicians who treat seriously impaired and disabled children. People would come to distrust physicians, which in turn would damage the entire health care system.

IN DEPTH UTILITARIANISM AND THE GOLDEN RULE

Probably much to the dismay of his religious critics, John Stuart Mill defended his radical doctrine of utilitarianism by arguing that it was entirely consistent with a fundamental Christian teaching:

In the golden rule of Jesus of Nazareth, we read the complete spirit of the ethics of utility. To do as one would be done by, and to love one's neighbour as oneself, constitute the ideal perfection of utilitarian morality. As the means of making the nearest approach to this ideal, utility would enjoin, first, that laws and social arrangements should place the happiness, or (as speaking practically it may be called) the

interest, of every individual, as nearly as possible in harmony with the interest of the whole; and secondly, that education and opinion, which have so vast a power over human character, should so use that power as to establish in the mind of every individual an indissoluble association between his own happiness and the good of the whole; especially between his own happiness and the practice of such modes of conduct, negative and positive, as regard for the universal happiness prescribes: so that not only he may be unable to conceive the possibility of happiness to himself, consistently with conduct opposed to the general good, but also that a direct impulse to promote the general good may be in every individual one of the habitual motives of action, and the sentiments connected therewith may fill a large and prominent place in every human being's sentient existence.⁴

Society might begin to devalue the lives of disabled people generally as well as the elderly and other vulnerable populations. The rule would also appear to entail a blatant violation of the cardinal principle of medical practice—do no harm. Adhering to it might therefore cause an erosion of all ethical codes and professional standards in medicine. But following the first rule would have no such consequences. It would permit the suffering of some impaired children, but this consequence seems not to be as catastrophic as those produced by consistently conforming to the second rule. For the rule-utilitarian, then, the morally right action would be *not* to administer the lethal injection, despite the parents' pleas.

Kantian Ethics

From the great German philosopher Immanuel Kant (1724–1804) comes what is widely regarded as probably the most sophisticated and influential deontological theory ever devised. It is the very antithesis of utilitarianism, holding that right actions do not depend in the least on consequences, the maximization of utility, the

production of happiness, or the desires and needs of human beings. For Kant, the core of morality consists of following a rational and universally applicable moral rule and doing so solely out of a sense of duty. An action is right only if it conforms to such a rule, and we are morally praiseworthy only if we perform it for duty's sake alone.

In Kant's system, all our moral duties are expressed in the form of *categorical imperatives*. An imperative is a command to do something; it is categorical if it applies without exception and without regard for particular needs or purposes. A categorical imperative says, "Do this—regardless." In contrast, a *hypothetical imperative* is a command to do something if we want to achieve particular aims, as in "If you want good pay, work hard." The moral law, then, rests on absolute directives that do not depend on the contingencies of desire or utility.

Kant says that through reason and reflection we can derive our duties from a single moral principle, what he calls *the categorical imperative*. He formulates it in different ways, the first one being "Act only on that maxim through

which you can at the same time will that it should become a universal law."⁵ For Kant, our actions have logical implications—they imply general rules, or maxims, of conduct. If you tell a lie for financial gain, you are in effect acting according to a maxim like "It's okay to lie to someone when doing so benefits you financially." The question is whether the maxim corresponding to an action is a legitimate moral law. To find out, we must ask if we could consistently will that the maxim become a universal law applicable to everyone—that is, if everyone could consistently act on the maxim and we would be willing to have them do so. If we could do this, then the action described by the maxim is morally permissible; if not, it is prohibited. Thus moral laws embody two characteristics thought to be essential to morality itself: universality and impartiality.

To show us how to apply this formulation of the categorical imperative to a specific situation, Kant uses the example of a lying promise. Suppose you need to borrow money from a friend, but you know you could never pay her back. So to get the loan, you decide to lie, falsely promising to repay the money. To find out if such a lying promise is morally permissible, Kant would have you ask if you could consistently will the maxim of your action to become a universal law, to ask, in effect, "What would happen if everyone did this?" The maxim is "Whenever you need to borrow money you cannot pay back, make a lying promise to repay." So what *would* happen if everyone in need of a loan acted in accordance with this maxim? People would make lying promises to obtain loans, but everyone would also know that such promises were worthless, and the custom of loaning money on promises would disappear. So willing the maxim to be a universal law involves a contradiction: If everyone made lying promises, promise-making itself would be no more; you cannot consistently will the maxim to become a universal law. Therefore, your duty is clear: Making a lying promise to borrow money is morally wrong.

Kant's first formulation of the categorical imperative yields several other duties, some of which

are particularly relevant to bioethics. Notably he argues that there is an absolute moral prohibition against killing the innocent, lying, committing suicide, and failing to help others when feasible.

Perhaps the most renowned formulation of the categorical imperative is the principle of respect for persons (a formulation distinct from the first one, though Kant thought them equivalent). As he puts it, "Act in such a way that you always treat humanity, whether in your own person or in the person of any other, never simply as a means, but always at the same time as an end."⁶ People must never be treated as if they were mere instruments for achieving some further end, for people are ends in themselves, possessors of ultimate inherent worth. People have ultimate value because they are the ultimate source of value for other things. They bestow value; they do not have it bestowed upon them. So we should treat both ourselves and other persons with the respect that all inherently valuable beings deserve.

According to Kant, the inherent worth of persons derives from their nature as free, rational beings capable of directing their own lives, determining their own ends, and decreeing their own rules by which to live. Thus, the inherent value of persons does not depend in any way on their social status, wealth, talent, race, or culture. Moreover, inherent value is something that all persons possess equally. Each person deserves the same measure of respect as any other.

Kant explains that we treat people merely as a means instead of an end-in-themselves if we disregard these characteristics of personhood—if we thwart people's freely chosen actions by coercing them, undermine their rational decision-making by lying to them, or discount their equality by discriminating against them. In bioethics, clear-cut cases of not respecting persons in Kant's sense would normally include experimenting on people without their knowledge and consent, lying to them about their medical condition and prognosis, and forcing patients to receive treatment against their will.

Notice that this formulation of the categorical imperative does not actually prohibit treating a person as a means but forbids treating a person *simply*, or *merely*, as a means—as nothing but a means. Kant recognizes that in daily life we often must use people to achieve our various ends. To buy milk, we use the cashier; to find books, we use the librarian; to get well, we use the doctor. But because their actions are freely chosen and we do not undermine their status as persons, we do not use them *solely* as instruments of our will. Medical researchers use their human subjects as a means to an end—but not merely as a means to an end if the subjects give their informed consent to participate in the research.

Natural Law Theory

From ancient times to the present day, many people have thought that the outlines of the moral law are plain to see because they are written large and true in nature itself. This basic notion has been developed over the centuries into what is known as **natural law theory**, the view that right actions are those that conform to moral standards discerned in nature through human reason. Undergirding this doctrine is the belief that all of nature (including humankind) is teleological, that it is somehow directed toward particular goals or ends, and that humans achieve their highest good when they follow their true, natural inclinations leading to these goals or ends. There is, in other words, a way things *are*—natural processes and functions that accord with the natural law—and how things are shows how things *should be*. The prime duty of humans, then, is to guide their lives toward these natural ends, acting in accordance with the requirements of natural law.

Implicit in all this is the element of rationality. According to natural law theory, humans are rational beings empowered by reason to perceive the workings of nature, determine the natural inclinations of humans, and recognize the implications therein for morally permissible actions. That is, reason enables human beings to ascertain the moral law implicit in nature and to apply that objective, universal standard to their lives.

Though natural law theory has both religious and nonreligious forms, the theistic formulation of theologian-philosopher Thomas Aquinas (1225–1274) has been the theory's dominant version. It is not only the official moral outlook of the Roman Catholic Church, but it has also been the intellectual starting point for many contemporary variations of the theory, secular and otherwise. For Aquinas, God is the author of the natural law who gave humans the gift of reason to discern the law for themselves and live accordingly. Aquinas argues that human beings naturally tend toward—and therefore have a duty of—preserving human life and health (and so must not kill the innocent), producing and raising children, seeking knowledge (including knowledge of God), and cultivating cooperative social relationships. In all this, Aquinas says, the overarching aim is to do and promote good and avoid evil.

Natural law theory does not provide a relevant moral rule covering every situation, but it does offer guidance through general moral principles, some of which are thought to apply universally and absolutely (admitting no exceptions). Among these principles are absolutist prohibitions against directly killing the innocent, lying, and using contraceptives. In his list of acts considered wrong no matter what, Aquinas includes adultery, blasphemy, and sodomy.

Of course, moral principles or rules often conflict, demanding that we fulfill two or more incompatible duties. We may be forced, for example, to either tell a lie and save people's lives or tell the truth and cause their death—but we cannot do both. Some moral theories address these problems by saying that all duties are *prima facie*: When duties conflict, we must decide which ones override the others. Theories that posit absolute duties—natural law theory being a prime example—often do not have this option. How does the natural law tradition resolve such dilemmas? Among other resources, it uses the **doctrine of double effect**.

This principle, a cornerstone of Roman Catholic ethics, affirms that performing a bad action to bring about a good effect is never

morally acceptable but that performing a good action may sometimes be acceptable even if it produces a bad effect. More precisely, the principle says it is always wrong to intentionally perform a bad action to produce a good effect, but doing a good action that results in a bad effect may be permissible if the bad effect is not intended although foreseen. In the former case, a bad thing is said to be directly intended; in the latter, a bad thing is not directly intended.

These requirements have been detailed in four "tests" that an action must pass to be judged morally permissible. We can express a traditional version of these tests like this:

1. The action itself must be morally permissible.
2. Causing a bad effect must not be used to obtain a good effect (the end does not justify the means).
3. Whatever the outcome of an action, the intention must be to cause only a good effect (the bad effect can be foreseen but never intended).
4. The bad effect of an action must not be greater in importance than the good effect.

Consider the application of these tests to euthanasia. Suppose an 80-year-old hopelessly ill patient is in continuous, unbearable pain and begs to be put out of her misery. Is it morally permissible to grant her request (either by giving a lethal injection or ending all ordinary life-sustaining measures)? If we apply the doctrine of double effect as just outlined, we must conclude that the answer is *no*: Euthanasia—either active or passive—is not a morally permissible option here. (In the Roman Catholic view, all forms of euthanasia are wrong, although it is permissible not to treat a hopelessly ill person for whom ordinary life-sustaining treatments are useless.) Failing even one of the tests would render an action impermissible, but in this case let us run through all four as a natural law theorist might:

1. Taking steps to terminate someone's life is a clear violation of test 1. Whatever its

effects, the action of taking a life is in itself immoral, a violation of the cardinal duty to preserve innocent life.

2. Ending the woman's life to save her from terrible suffering is an instance of causing a bad effect (the woman's death) as a means of achieving a good effect (cessation of pain)—a failure of test 2.
3. The death of the woman is intended; it is not merely a tragic side effect of the attempt solely to ease her pain. So the action fails test 3.
4. Causing the death of an innocent person is a great evil that cannot be counterbalanced by the good of pain relief. So the action does not pass test 4.

The verdict in such a case would be different, however, if the patient's death were not intentionally caused but unintentionally brought about. Suppose, for example, that the physician sees that the woman is in agony and so gives her a large injection of morphine to minimize her suffering—knowing full well that the dose will also probably speed her death. In this scenario, the act of easing the woman's pain is itself morally permissible (test 1). Her death is not a means to achieve some greater good; the goal is to ease her suffering (test 2). Her death is not intended; the intention is to alleviate her pain, though the unintended (but foreseen) side effect is her hastened death (test 3). Finally, the good effect of an easier death seems more or less equivalent in importance to the bad effect of a hastened death. Therefore, unintentionally but knowingly bringing about the woman's death in this way is morally permissible.

We get similar results if we apply the double-effect principle in the traditional way to abortion. We find that as the intentional destruction of an innocent human life (so-called direct), abortion is always immoral (test 1). Moreover, it is wrong even (or especially) if it is performed to bring about some good result, such as saving the mother's life or preventing serious harm to her (tests 2 and 3). On the other hand, actions

leading unintentionally to the death of a fetus (so-called indirect abortion) may be permissible in rare cases. Say a pregnant woman has an infectious disease that will kill her unless she gets injections of a powerful drug. But the drug will abort the fetus. According to the doctrine of double effect, receiving the injections may be morally permissible if the action itself is morally permissible, which it is (test 1); if the death of the fetus is not used to rescue the woman (test 2); if the injections are given with the intention of curing the woman's disease, not of inducing an abortion (test 3); and if the death of the fetus is balanced by the life of the woman (test 4).

Rawls' Contract Theory

In its broadest sense, **contractarianism** refers to moral theories based on the idea of a social contract, or agreement, among individuals for mutual advantage. The most influential contemporary form of contractarianism is that of philosopher John Rawls (1921–2002), who uses the notion of a social contract to generate and defend moral principles governing how members of a society should treat one another. He asks, in effect, by what principles should a just society structure itself to ensure a fair distribution of rights, duties, and advantages of social cooperation?

His answer is that the required principles—essentially principles of justice—are those that people would agree to under hypothetical conditions that ensure fair and unbiased choices. He believes that if the starting point for the social contract is fair—if the initial conditions and bargaining process for producing the principles are fair—then the principles themselves will be just and will define the essential makeup of a just society. As Rawls says,

[The] guiding idea is that the principles of justice for the basic structure of society are the object of the original agreement. They are the principles that free and rational persons concerned to further their own interests would accept in an initial position of equality as defining the fundamental terms of their association. These

principles are to regulate all further agreements; they specify the kinds of social cooperation that can be entered into and the forms of government that can be established.⁷

At the hypothetical starting point—what Rawls calls the “original position”—a group of normal, self-interested, rational individuals come together to choose the principles that will determine their basic rights and duties and their share of society's benefits and burdens. But to ensure that their decisions are as fair and impartial as possible, they must meet behind a metaphorical “veil of ignorance.” Behind the veil, no one knows his own social or economic status, class, race, sex, abilities, talents, level of intelligence, or psychological makeup. Since the participants are rational and self-interested but ignorant of their situation in society, they will not agree to principles that will put any particular group at a disadvantage because they might very well be members of that group. They will choose principles that are unbiased and nondiscriminatory. The assumption is that since the negotiating conditions in the original position are fair, the agreements reached will also be fair—the principles will be just.

Rawls contends that given the original position, the participants would agree to arrange their social relationships according to these fundamental principles:

1. Each person is to have an equal right to the most extensive total system of equal basic liberties compatible with a similar system of liberty for all.
2. Social and economic inequalities are to be arranged so that they are both:
 - (a) to the greatest benefit of the least advantaged . . . and
 - (b) attached to offices and positions open to all under conditions of fair equality of opportunity.⁸

The first principle—the equal liberty principle—says that everyone is entitled to the most freedom possible in exercising basic rights

and duties (for example, the right to vote and hold office and freedom of speech, assembly, and thought). Each person should get a maximum degree of basic liberties but no more than anyone else. This principle takes precedence over all other considerations (including the second principle) so that basic liberties cannot be reduced or canceled just to improve economic well-being.

The second principle concerns social and economic goods such as income, wealth, opportunities, and positions of authority. Part (b) says that everyone is entitled to an equal chance to try to acquire these basic goods. No one is guaranteed an equal share of them, but opportunities to obtain these benefits must be open to all, regardless of social standing.

Rawls knows that social and economic inequalities will naturally arise in society. But as he asserts in part (a), they are not unjust if they work to everyone's benefit, especially to the benefit of the least well off in society. "[There] is no injustice," he says, "in the greater benefits earned by a few provided that the situation of persons not so fortunate is thereby improved."⁹ For Rawls, such a policy is far more just than one in which some people are made to suffer for the greater good of others: "[I]t is not just that some should have less in order that others may prosper."

In Rawls' scheme, the demands of the first principle must be satisfied before satisfying the second, and the requirements of part (b) must be met before those of part (a). In any just distribution of benefits and burdens, then, the first priority is to ensure equal basic liberties for all concerned, then equality of opportunity, then the arrangement of any inequalities to the benefit of the least advantaged.

As a theory of distributive justice, Rawls' contractarianism seems to have significant implications for the allocation of society's health care resources. For example, one prominent line of argument goes like this: As Rawls claims, everyone is entitled to fair equality of opportunity, and adequate (basic) health care enables fair equality of opportunity (by ensuring "normal species functioning"). Therefore, everyone is

entitled to adequate health care, and all appropriate measures compensating for the disadvantage and impairment.¹⁰ In such a system, there would be universal access to a basic level of health care, while more elaborate or elective services would be available to anyone who could afford them.

Another implication: Suppose that to provide a basic level of health care to everyone (and meet the equality-of-opportunity requirement), society would have to spend 90 percent of its health care resources. But say that in the current system, 50 percent of the resources are being spent on acute care for the elderly—that is, expensive measures to extend the lives of people who have already lived a long time. According to Rawlsian principles, is the current system of health care unjust?

Virtue Ethics

Most moral theories—including all those just discussed—are theories of obligation. They emphasize the rightness of actions and the duties of moral agents. Their main concern is knowing and doing what's right, and their chief guide to these aims is moral principles or directives. Virtue ethics, however, is a radically different kind of moral theory: It focuses on the development of virtuous character. According to virtue ethics, character is the key to the moral life, for it is from a virtuous character that moral conduct and values naturally arise. Virtues are ingrained dispositions to act by standards of excellence, so having the proper virtues leads as a matter of course to right actions properly motivated. The central task in morality, then, is not knowing and applying principles but being and becoming a good person, someone possessing the virtues that define moral excellence. In virtue ethics, someone determines right action not by consulting rules but by asking what a truly virtuous person would do or whether an action would accord with the relevant virtues.

Aristotle (384–322 B.C.E.) is the primary inspiration for contemporary versions of virtue ethics. For him, as for many modern virtue

and duties (for example, the right to vote and hold office and freedom of speech, assembly, and thought). Each person should get a maximum degree of basic liberties but no more than anyone else. This principle takes precedence over all other considerations (including the second principle) so that basic liberties cannot be reduced or canceled just to improve economic well-being.

The second principle concerns social and economic goods such as income, wealth, opportunities, and positions of authority. Part (b) says that everyone is entitled to an equal chance to try to acquire these basic goods. No one is guaranteed an equal share of them, but opportunities to obtain these benefits must be open to all, regardless of social standing.

Rawls knows that social and economic inequalities will naturally arise in society. But as he asserts in part (a), they are not unjust if they work to everyone's benefit, especially to the benefit of the least well off in society. "[There] is no injustice," he says, "in the greater benefits earned by a few provided that the situation of persons not so fortunate is thereby improved."⁹ For Rawls, such a policy is far more just than one in which some people are made to suffer for the greater good of others: "[I]t is not just that some should have less in order that others may prosper."

In Rawls' scheme, the demands of the first principle must be satisfied before satisfying the second, and the requirements of part (b) must be met before those of part (a). In any just distribution of benefits and burdens, then, the first priority is to ensure equal basic liberties for all concerned, then equality of opportunity, then the arrangement of any inequalities to the benefit of the least advantaged.

As a theory of distributive justice, Rawls' contractarianism seems to have significant implications for the allocation of society's health care resources. For example, one prominent line of argument goes like this: As Rawls claims, everyone is entitled to fair equality of opportunity, and adequate (basic) health care enables fair equality of opportunity (by ensuring "normal species functioning"). Therefore, everyone is

entitled to adequate health care, which includes all appropriate measures for eliminating or compensating for the disadvantages of disease and impairment.¹⁰ In such a system, there would be universal access to a basic level of health care, while more elaborate or elective services would be available to anyone who could afford them.

Another implication: Suppose that to provide a basic level of health care to everyone (and meet the equality-of-opportunity requirement), society would have to spend 90 percent of its health care resources. But say that in the current system, 50 percent of the resources are being spent on acute care for the elderly—that is, expensive measures to extend the lives of people who have already lived a long time. According to Rawlsian principles, is the current system of health care unjust?

Virtue Ethics

Most moral theories—including all those just discussed—are theories of obligation. They emphasize the rightness of actions and the duties of moral agents. Their main concern is knowing and doing what's right, and their chief guide to these aims is moral principles or directives. Virtue ethics, however, is a radically different kind of moral theory: It focuses on the development of virtuous character. According to virtue ethics, character is the key to the moral life, for it is from a virtuous character that moral conduct and values naturally arise. Virtues are ingrained dispositions to act by standards of excellence, so having the proper virtues leads as a matter of course to right actions properly motivated. The central task in morality, then, is not knowing and applying principles but being and becoming a good person, someone possessing the virtues that define moral excellence. In virtue ethics, someone determines right action not by consulting rules but by asking what a truly virtuous person would do or whether an action would accord with the relevant virtues.

Aristotle (384–322 B.C.E.) is the primary inspiration for contemporary versions of virtue ethics. For him, as for many modern virtue

ethicists, the highest goal of humanity is the good life, or "human flourishing" (what Aristotle calls *eudaimonia*, or happiness), and developing virtues is the way to achieve such a rich and satisfying life. Thus virtues are both the traits that make us good persons and the dispositions that enable us to live good lives. The good life is the virtuous life.

Unlike many theories of obligation, virtue ethics asks us to do more than just observe minimal moral rules—it insists that we *aspire to moral excellence*, that we cultivate the virtues that will make us better persons. In this sense, virtue ethics is goal-directed, not rule-guided. The moral virtues—benevolence, honesty, loyalty, compassion, fairness, and the like—are ideals that we must ever strive to attain. (There are also nonmoral virtues such as patience, prudence, and reasonableness, which need not concern us here.) By the lights of both Aristotle and modern virtue ethicists, character is not static. We can become more virtuous by reflecting on our lives and those of others, practicing virtuous behavior, or imitating moral exemplars such as Gandhi, Buddha, Jesus, Muhammad, and Socrates. We can—and should—be better than we are.

To the virtue ethicist, possessing the right virtues means having the proper motivations that naturally accompany those virtues. To act morally, we must act from virtue, and acting from virtue means acting with the appropriate motives. It is not enough to do right; we must do right for the right motivating reasons. If we save a drowning friend, we should do so out of genuine feelings of compassion, kindness, or loyalty—not because of the prodding of moral rules or social expectations. In contrast, some moral theories (notably Kant's) maintain that acting morally is solely a matter of acting for duty's sake—performing an action simply because duty requires it. Virtuous motives are irrelevant; we act morally if we do our duty regardless of our motivations. But this notion seems to many to offer a barren picture of the moral life. Surely, they say, motivations for acting are often relevant to our evaluations of

people's character and actions. The friend we saved from drowning would probably be appalled if we declared that we saved her out of duty even though we did not really care whether she lived or died. Many moral philosophers agree that motivations are indeed important considerations in moral judgments, and they have incorporated virtues into their theories of obligation.

Virtue ethics fits well with the emphasis on virtues that has always been part of the healing arts. Physicians and nurses are expected to possess particular virtues, including compassion, trustworthiness, justice, and honesty. They are expected to be more than just technically skilled and knowledgeable and to do more than merely follow the rules of conduct or procedure. They are obliged to do right by their patients, and this obligation is most likely met through the cultivation and possession of virtues.

The virtue ethics approach to bioethical issues is distinctive. On abortion, for example, the virtue ethicist might argue that a woman's decision to have an abortion should be judged by the virtues (or lack thereof) that she draws on in deciding what to do. If she decides to have an abortion just because she is afraid of the responsibilities of parenthood, she shows cowardice. If she wants to go through with an abortion merely because pregnancy would disrupt her vacation plans, she shows self-centeredness and callousness. In neither case is the virtue ethicist likely to call the woman's decision virtuous.¹¹

The Ethics of Care

The ethics of care is a distinctive moral perspective that arose out of feminist concerns and grew to challenge core elements of most other moral theories. Generally those theories emphasize abstract principles, general duties, individual rights, impartial judgments, and deliberative reasoning. But the ethics of care shifts the focus to the unique demands of specific situations and to the virtues and feelings that are central to close personal relationships—empathy, compassion, love, sympathy, and fidelity. The heart of the moral life is feeling for and caring for those with